

Enterprise Architecture Design — Supplemental Assignment

Module: ENTPH6001 Enterprise Architecture Design **Submission:** Cloud-native architecture for a small e-commerce checkout platform **Word count (Sections 1–5, excluding figures, references and appendices):** 1996

1. Architecture Proposal

The organisation runs its checkout application as a small set of containers, which leaves every dependency implicit: anything can call anything, configuration travels with the image, and one slow component stalls the customer path. The proposal makes those relationships enforceable.

Service decomposition. The system is decomposed along business capability rather than technical layer, following bounded-context reasoning: `pricing-fn` owns tax and totals, `inventory-fn` owns stock and the stock datastore, and `checkout-fn` composes them into an order decision. A separate `gateway` serves the UI and exposes one API route. The split is deliberately small: at this size the dominant risk is operational overhead, not team scaling, so services were justified by differing rates of change and failure isolation rather than by count. `inventory-fn` alone owns its schema; no other service reaches the database, which prevents the shared-database coupling that turns a distributed system back into a distributed monolith.

Request entry. External traffic enters through one controlled point: a Traefik Ingress routing to `gateway-svc`. The gateway terminates the public contract, exposes only `/api/checkout`, and either honours an inbound `X-Request-ID` or mints one. Internal services publish no external route. **Trust tiers.** Workloads are classified **public** (gateway), **internal** (checkout, pricing, inventory) and **protected** (PostgreSQL). The classification is not documentation: a deny-by-default NetworkPolicy covering **both ingress and egress** (Kubernetes, 2026b), plus directed allow-rules, makes each tier boundary a control the platform enforces. Omitting egress rules is a serious gap: a compromised workload with unrestricted outbound access can exfiltrate data or reach a cloud metadata endpoint, the mechanism behind the Capital One breach (Lee, 2026). Here `pricing-fn` and `postgres` are declared with empty egress, since neither has reason to initiate a connection, and only `inventory-fn` may egress to the database.

State, configuration and secrets. Stock data lives in PostgreSQL backed by a PersistentVolumeClaim. Non-secret configuration (dependency URLs, timeout budget, tax rate, database host) comes from a ConfigMap; credentials come from a Secret injected as environment variables. Nothing is baked into an image, and the repository holds only a Secret *template*, the real value being generated at deploy time.


Workload identity. Each workload runs under its own ServiceAccount holding no API permissions, with token automounting disabled. Sharing the namespace `default` account makes API activity unattributable and grants access by accident; separate identities cost nothing and make any future grant precise.

Migration path. Because the gateway owns the public contract, routes can move behind it one at a time — the strangler-fig pattern (Fowler, 2004) — so the existing deployment keeps serving while capabilities are cut over. That is why the gateway is a distinct workload, not ingress annotations.

Operational and security boundaries. Two decisions carry most of the resilience argument. First, `checkout-fn` applies a 1500 ms timeout budget per dependency, so a slow dependency cannot hold customer requests open. Second, dependencies are classified: pricing is **critical** (no total, no order), inventory **non-critical** (the order can be accepted with stock unconfirmed). Health probes follow the same logic (Kubernetes, 2026a). Liveness is shallow everywhere, because a liveness probe that touches a dependency

converts a dependency outage into a restart storm; only `inventory-fn`, which owns the datastore, probes it in readiness. `checkout-fn` readiness excludes its dependencies, since probing them would remove every replica during a dependency outage and convert partial failure into total outage.

Figure 1 — Implemented architecture with trust tiers, workload identity and enforced policy. Solid arrows are permitted paths; dotted DENIED / EGRESS DENIED edges are refused by NetworkPolicy and are the cases verified in Appendix D.



2. Implementation Summary

The prototype runs on single-node K3s (v1.36.3) on Ubuntu 24.04 (Rancher, 2026), installed with the module's procedure. Images are built locally and imported into containerd, so no registry is assumed.


Implemented: four Deployments (gateway ×2, checkout ×2, pricing ×2, inventory ×2) plus PostgreSQL; five ClusterIP Services; one Ingress carrying an edge rate-limit middleware; two ConfigMaps; one Secret; one PVC; twelve NetworkPolicies covering both directions; probes, resource limits and `restricted` Pod Security (Kubernetes, 2026c) on every workload, with non-root execution, dropped capabilities and no token mounting.

Four choices map design intent onto configuration. Two replicas everywhere make the readiness gate meaningful: a bad replica can be withheld while another serves. `TIMEOUT_MS=1500` bounds a slow dependency. PostgreSQL uses `Recreate`, not rolling update, because one RWO volume cannot be shared by two pods mid-rollout. The gateway runs `nginx-unprivileged` because stock NGINX cannot satisfy the `restricted` standard.

Structured JSON logging with request-ID propagation was added to the three application services. The platform deploys with one command and was verified by deleting the namespace and rebuilding: teardown 61 seconds, rebuild to a working checkout **30 seconds** (Appendix E).

Gaps against the proposal. Three things are designed but not implemented, for scope not oversight: the Ingress serves HTTP (no certificate authority in scope), the database is a single replica with no backup or failover, and the checkout API has no authentication. Each is revisited in Section 5.

Figure 2 — Main request flow and the evidence point behind each claim.



3. Operational Behaviour

The behaviour selected is **partial failure: graceful degradation versus fail fast**, because it tests the dependency classification directly. Healthy, a checkout returns `200 confirmed` in 10 ms. With `inventory-fn` scaled to zero it returns 200, marked degraded, in 1502 ms. With `pricing-fn` delayed 3000 ms it returns `503 pricing_timeout` in 1504 ms (Appendix B).

Both claims hold: the non-critical dependency degrades rather than rejecting the customer, and the critical one fails fast at the configured budget. The measured 1502/1504 ms confirms the 1500 ms timeout is the binding constraint, not an accident of network timing. The cost is visible too — degradation turns a 10 ms response into a 1502 ms one, because the timeout must expire first. Under sustained failure that consumes connection capacity (Google, 2016); a circuit breaker is the correct next step (Fowler, 2014).

A second experiment came from a real defect: an intermittent 502 during a rollout, reproduced by the control group. A `preStop` delay plus a 30-second grace period lets the endpoints controller remove the pod first (Kubernetes, 2026e). Run as a controlled A/B rather than a single run, the control failed 6 of 179 requests (3.4%) and the fixed group 3 of 208 (1.4%). At these sample sizes that difference is not statistically significant, and the control also ran a 1-second grace period against the 30-second default, so two variables moved together. What it does establish is that failures were **not eliminated**, and that an earlier run returning 204/204 was unrepresentative.

A third experiment tests admission control, which is load shedding at the edge. A 30-request burst without the rate-limit middleware (Traefik Labs, 2026) admitted all 30 at a p95 of 0.168 s; with it, 11 were admitted and 19 refused with HTTP 429, and the admitted requests completed at a p95 of 0.075 s. The two samples differ in size and selection, so this shows the direction of the effect rather than its magnitude: admitting everything degrades everyone, whereas early rejection keeps the backend out of contention.

A fourth set asked what happens when a *change* is wrong rather than a dependency slow. A broken readiness path, a broken image tag, and the gateway upstream reverted to its Compose name were injected separately (Appendix C3–C5). All three were contained by one mechanism — a replica receives traffic only after passing its readiness gate — and user-visible impact was nil: 10/10, 15/15 and 5/5 checkouts returned 200. Operators saw the failure; users did not, so a broken release can sit unnoticed unless rollout state is alerted on.

4. Observability and Security Validation

Observability. The three application services emit structured JSON with a shared `request_id`, and `checkout-fn` adds per-dependency `duration_ms`. One request is traceable across all three, and the per-hop timings identified the timeout as the binding constraint in Section 3. This is deliberately log-based rather than a full tracing stack: on a single node, correlation IDs deliver most of the diagnostic value far more cheaply. The limitation is real — no sampling, retention, aggregation or dashboard. **Security validation 1 — are the trust boundaries real?** The tiers were tested, not asserted: four documented paths must work, six lateral paths must fail, five outbound paths must fail, and an unlabelled pod simulating a compromised workload must reach nothing. **All 19 cases passed** (Appendix D). The egress cases matter most: no application workload can reach the public internet, and the cloud metadata endpoint at 169.254.169.254 is unreachable from the service holding database credentials. NetworkPolicy objects are silently inert without a policy-capable CNI, so the negative controls prove enforcement rather than intent. DNS is the one deliberate exception; without it service discovery fails and the outage looks like an application bug.

The limits matter as much as the result. The 19 cases cover four workloads, not Postgres; traffic to and from the hosting node is always permitted regardless of policy; and NetworkPolicy records nothing about what it allowed or blocked, so this control has no audit trail. It isolates reachability, and nothing more. **Security validation 2 — configuration and image posture.** `kubesecc` (Controlplane, 2026) scores all five Deployments, read from the live cluster so the score describes what actually runs. The four application workloads score **12** (a raw score, not a ratio) — non-root, dropped capabilities, read-only root filesystem, resource limits, dedicated identities, no token automounting. Postgres scores **11**, missing only `ReadOnlyRootFilesystem`. Trivy (Aqua Security, 2026) reported 1 CRITICAL and 6 HIGH per service image, including CVE-2026-59873 in `node-tar`. Interpretation mattered more than the count: every finding came from the npm CLI in the base image, not application dependencies, and was unreachable at runtime. Rather than suppress them, the images were rebuilt multi-stage without the package manager; the same scan now returns **0 CRITICAL and 0 HIGH** (Appendix F).

Prioritised issues. (1) Secrets are base64 only and K3s disables encryption at rest (Kubernetes, 2026d) — any principal with `get secret` reads the database password; this is the highest residual risk. (2) No authentication on the checkout API. (3) No TLS. (4) Lower priority: no AppArmor or seccomp profile and UIDs below 10000, flagged by `kubesecc` but minor given deny-by-default networking, dedicated identities and disabled token mounting.

5. Critical Evaluation

The strongest aspect is that the architecture's claims are testable and were tested: tier separation, degradation, timeout budget, persistence and rollout safety each have a script producing evidence. One script found a real defect; the control group used to measure the fix also revealed the fix is partial.

The weaknesses are equally clear. **Single points of failure:** one node, one database replica, so node loss is total outage — RPO and RTO are undefined and backups do not exist. **Identity is missing:** NetworkPolicy enforces reachability, not identity, so the trust model is entirely network-based, contradicting the Zero Trust principle that network location must not imply trust (NIST, 2020); a compromised `checkout-fn` is indistinguishable from a legitimate one to `inventory-fn`. mTLS with workload identity is the most valuable next increment. **Degradation is timeout-bound, not circuit-broken. The PVC gives restart survival, not recoverability** — it outlived pod replacement but was destroyed with the namespace during the rebuild test.

Two trade-offs were accepted deliberately. A service mesh would supply mTLS, retries and telemetry uniformly, but for four services it adds a control plane and sidecar overhead this organisation cannot operate; NetworkPolicy plus application timeouts covers most of the benefit far more cheaply. Scale-to-zero was rejected because cold starts would add latency to the checkout path and idle cost is not the constraint at this size. A third alternative, asynchronous messaging, is not settled. The two dependencies are called in parallel, so tail latency is bounded by the slower rather than their sum; what is still coupled is availability. The degraded response returned when inventory is unavailable is in effect a promise to confirm stock later — and nothing records or honours it. A durable queue with idempotent consumers and a dead-letter path is the standard remedy, rejected here only because a stateful broker would add a second single point of failure on one node. The promise remains outstanding, so this is an acknowledged gap, not a resolved trade-off.

If continued: encryption at rest and a secrets manager; TLS and authentication at the edge; workload identity for east-west calls; a circuit breaker on inventory; then backup and restore rehearsal — recovery that has never been tested is an assumption, not a control.

6. References

Aqua Security (2026) *Trivy documentation*. Available at: <https://trivy.dev/> (Accessed: 14 August 2026).

Controlplane (2026) *kubesecc — security risk analysis for Kubernetes resources*. Available at: <https://kubesecc.io/> (Accessed: 14 August 2026).

Fowler, M. (2004) *StranglerFigApplication*. Available at: <https://martinfowler.com/bliki/StranglerFigApplication.html> (Accessed: 14 August 2026).

Fowler, M. (2014) *CircuitBreaker*. Available at: <https://martinfowler.com/bliki/CircuitBreaker.html> (Accessed: 14 August 2026).

Google (2016) *Site Reliability Engineering: Addressing Cascading Failures*. Available at: <https://sre.google/sre-book/addressing-cascading-failures/> (Accessed: 14 August 2026). Kubernetes (2026a) *Configure Liveness, Readiness and Startup Probes*. Available at: <https://kubernetes.io/docs/tasks/configure-pod-container/configure-liveness-readiness-startup-probes/> (Accessed: 14 August 2026).

Kubernetes (2026b) *Network Policies*. Available at: <https://kubernetes.io/docs/concepts/services-networking/network-policies/> (Accessed: 14 August 2026).

Kubernetes (2026c) *Pod Security Standards*. Available at: <https://kubernetes.io/docs/concepts/security/pod-security-standards/> (Accessed: 14 August 2026).

Kubernetes (2026d) *Encrypting Confidential Data at Rest*. Available at: <https://kubernetes.io/docs/tasks/administer-cluster/encrypt-data/> (Accessed: 14 August 2026).

Kubernetes (2026e) *Pod Lifecycle: termination of Pods*. Available at: <https://kubernetes.io/docs/concepts/workloads/pods/pod-lifecycle/> (Accessed: 14 August 2026).

Lee, E. (2026) *Enterprise Architecture Design: Lectures 1–12 and Laboratory Materials*. ENTPH6001, TU Dublin.

NIST (2020) *SP 800-207: Zero Trust Architecture*. doi: 10.6028/NIST.SP.800-207.

Rancher (2026) *K3s documentation: Networking and Traefik Ingress*. Available at: <https://docs.k3s.io/networking> (Accessed: 14 August 2026).

Traefik Labs (2026) *Traefik Kubernetes Ingress documentation*. Available at: <https://doc.traefik.io/traefik/providers/kubernetes-ingress/> (Accessed: 14 August 2026).

AI tools declaration

GitHub Copilot (Claude Opus 5) was used to assist with drafting manifests, service code, test scripts and report structure. All deployments, experiments and measurements reported here were executed by the author on the cluster described, and every figure quoted is taken from the captured evidence logs in the appendices. No result in this report is estimated or reproduced from memory.